

## Problem Set 3

**Note:** We will discuss the first problem in the problem-solving session. However, you still need to write your own solution to every problem.

**Problem 1:** Learning to maximize two-dimensional functions ..... (10 points)

Consider the following maximization problem. In each round  $i = 1, 2, \dots, n$  you pick a pair of numbers  $(x_1^i, x_2^i) \in [0, 1]^2$ . In round  $i$ , nature picks an arbitrary function  $f^i : [0, 1]^2 \rightarrow [0, 1]$ , where  $f^i$ 's gradient is guaranteed to be in  $[-1, 1]^2$ . After the prediction is made, nature reveals only your payoff  $f^i(x_1^i, x_2^i)$ . You want to maximize your total payoff against the best-in-hindsight choice, i.e. minimizing the following regret:

$$R_n = \max_{(x_1, x_2) \in [0, 1]^2} \frac{1}{n} \sum_{i=1}^n f^i(x_1, x_2) - \frac{1}{n} \sum_{i=1}^n f^i(x_1^i, x_2^i).$$

Design an online learning algorithm for this problem and provide a bound for  $R_n$  of your algorithm. Clearly describe your algorithm and analyze it (state and derive your regret bound). The regret bound should only depend on  $n$  and vanish as  $n \rightarrow \infty$ . Explain your answer.

**Problem 2:** Learning to predict ..... (10 points)

Consider the following prediction problem. On each day  $t = 1, 2, \dots, n$  you predict the probability of raining  $p^t \in [0, 1]$ . After the prediction is made, nature reveals the state  $\theta^t \in \{0, 1\}$  (whether it really rains or not). You want to minimize your mean squared error on all days against the best-in-hindsight prediction, i.e. minimizing the following regret:

$$R_n = \frac{1}{n} \sum_{t=1}^n (p^t - \theta^t)^2 - \min_{p \in [0, 1]} \frac{1}{n} \sum_{t=1}^n (p - \theta^t)^2.$$

Design an online learning algorithm for this problem and provide a bound for  $R_n$  of your algorithm. Clearly describe your algorithm and analyze it (state and derive your regret bound). The regret bound should only depend on  $n$  and vanish as  $n \rightarrow \infty$ . Explain your answer.

**Hint:** Minimizing the regret for squared error is equivalent to minimizing the regret for reward ( $1 - \text{squared error}$ ). You might want to first discretize the prediction space to  $k$  levels, then use the EW algorithm and optimize  $k$ .